

A Brief Overview of the Actor-Critic Method
 A policy $\pi_\theta(A|S)$ gives the probability of taking action A given S .
 The goal then is to maximize the objective function

$$C_{t_0}(\theta) = E[\sum_{t=t_0}^{\infty} \log(\pi_\theta(a_t|s_t)) * r(a_t|s_t)],$$

where $r(A|S)$ represents the reward from taking action A at state S . We then define

$$Q(a_t|s_t) = C_t(\theta|a_t),$$

representing the q-value. The q-value of taking an action A at state S is equal to the expected sum of the discounted rewards, given that the action taken at time t is a_t . We now define two agents: an actor, $\pi_\theta(A|S)$, which predicts the probability of taking action A at state S , and critic $q_\theta(a_t|s_t)$, which predicts the value of $Q(a_t|s_t)$. Now, as the actor trains on the environment, the critic can be optimized by looking at the actor's total reward after an episode, and the actor is able to learn using the critic's predicted q-values. [TODO MAKE MORE DETAILED]